

Intel Cloud Orchestration Networking

Winter Final Report

Cody Malick

Team 51, Cloud Orchestra

Abstract

This document outlines the progress of the Cloud Orchestration Networking project over the Fall and Winter terms. It contains a short description of the project's purposes and goals, current progress, current issues, and any solutions to those issues. It also contains a week by week retrospective for all ten weeks of Fall term and the first half of Winter term.

CONTENTS

| | | |
|------------|-----------------------------|----------|
| I | Project Goals | 2 |
| II | Purpose | 2 |
| III | Current Progress | 2 |
| IV | Issues | 3 |
| IV-A | Environment Setup | 3 |
| V | Team Review | 5 |
| V-A | Garrett Smith | 5 |
| V-B | Matthew Johnson | 5 |
| VI | Retrospective | 5 |
| | References | 5 |

I. PROJECT GOALS

Our project is to first switch the Linux-created GRE tunnel implementation in Ciao to use GRE tunnels created by Open vSwitch. From that point we will switch the actual tunneling implementation from GRE to VxLAN/nvGRE based on performance measurements of each on data center networking cards. After this is completed, a stretch goal is to replace Linux bridges with Open vSwitch switch instances.

These goals changed somewhat by the middle of the Winter term. The primary goal now is to replace the Linux bridges with Open vSwitch switch instances because of an assumption that was found to be incorrect. This change will be discussed in the Winter Progress section. This change also pushes the other two goals back behind the bridge goal.

II. PURPOSE

The current implementation of Ciao tightly integrates software defined networking principles to leverage a limited local awareness of just enough of the global cloud's state. Tenant overlay networks are used to overcome traditional hardware networking challenges by using a distributed, stateless, self-configuring network topology running over dedicated network software appliances. This design is achieved using Linux-native Global Routing Encapsulation (GRE) tunnels and Linux bridges, and scales well in an environment of a few hundred nodes.

While this initial network implementation in Ciao satisfies current simple networking needs in Ciao, all innovation around software defined networks has shifted to the Open vSwitch (OVS) framework. Moving Ciao to OVS will allow leverage of packet acceleration frameworks like the Data Plane Development Kit (DPDK) as well as provide support for multiple tunneling protocols such as VxLAN and nvGRE. VxLAN and nvGRE are equal cost multipath routing (ECMP) friendly, which could increase network performance overall.

III. CURRENT PROGRESS

Early Winter term was defined by various errors created first by the operating system we were using, and further compounded by networking issues from the MAC locked OSU network. We worked through these issues one-by-one and managed to make progress with development of the first goal. This progress was stymied by the discovery that our assumption about Open vSwitch compatibility was wrong about halfway through the term, which has changed the scope of our project considerably. While we originally planned to have that first goal completed by the end of the fifth week, this scope change altered our schedule and pushed it back. Good progress has been made on testing the basic cluster setup, but this has been hindered by the OSU network restrictions.

In the latter half of the term, there was a large change in scope. We had discovered that, in order to utilize Open vSwitch, we had to use OVS generated bridges, as well as OVS generated tunnels. The primary difference between the old scope and the new is generating bridges using OVS. What this means is that there is quite a bit of implementation work that had to be done to get OVS bridges created. The changes were discussed and approved with the team at Intel. After further discussion, it was agreed that we could use a direct call to the OVS command line utility, `ovs-vsctl`.

To call the command line utility from Go code, we needed to use Go's `exec` function. What `exec` allows a program to do is call third party tools from a shell environment. Here is some sample code on creating an OVS generated bridge while passing a function's arguments as arguments to the command line tool:

Listing 1. Simple OVS bridge creation function

```
func createOvsBridge(bridgeId string) error {
    // Example: ovs-vsctl add-br ovs-br1
    args := []string{"add-br", bridgeId}
    cmd, err := exec.Command("ovs-vsctl", args...).Start()
    if err != nil {
        return err
    }
    return nil
}
```

With functions similar to this one, implementation of the OVS command line utility has gone quite smoothly. The main portion of the work that has been done in the last few weeks have been focused on adding a new network mode to Ciao. With the addition of a new type of bridge, all the accompanying programmatic logic to change modes must be put in place. Ciao did not previously support anything outside of the default mode. As such, it has been quite challenging and fun figuring out the best way to implement our code into the existing architecture. We are well on our way to a prototype, and will continue work over spring break. Our goal to is to having a working alpha to show on returning next term.

IV. ISSUES

In this section, we will cover the issues we ran into at the beginning of term, as well as the solutions, and discussions that were had around them. The first two major issues are related to environment setup. These two issues consumed a large portion of the first part of the term.

A. Environment Setup

Environment setup is a pivotal part of our project, as developing and deploying to a non-cloud environment would not show what the project has accomplished. As such, it is important that our team have a good testing environment. The first, and arguably most major issue that was experienced, was setting up the Ciao cluster. The first step of the project is, quite simply, getting Ciao working in the state that it was provided to us. This ultimately should have been fairly straightforward, but we ran into two major issues. The first of which was the management of certificates in Clear Linux.

Clear Linux, Intel's custom Linux distribution, is a fast operating system designed to take advantage of Intel CPU's advanced features that often go under utilized. This was an obvious choice for our team. We initially set up all the NUCs with Clear Linux installed, and proceeded with deployment. We spent about a week tackling some minor configuration issues that came from simply being new to the system. Later, about a week was spent tracking down an issue within Clear Linux. Specifically,

that Clear Linux manually manages the network trust store. This isn't an issue in general, but this behavior of Clear Linux is slightly different than Ubuntu's. At first we thought we had simply set a value wrong in a configuration. But after further investigation, it became clear it was an issue unrelated to configuration. The difference in behavior was reported to the Clear Linux development team. They have since shipped a bug fix in a later release. On that front, we were able to iron out Ciao's behavior on Clear Linux.

The second major issue we encountered was that of getting FQDN's in Ciao, or Fully Qualified Domain Names. These were important to the controller node of Ciao as it needed to identify the compute nodes needed to deploy software onto. As we were trying to deploy, we found that Ciao kept attempting to connect to the hostname of the device instead of the FQDN. For example, the control node would try to connect to `fw-dear205-ciao-nuc0` instead of the FQDN, `fw-dear205-ciao-nuc0.engr.oregonstate.edu`. A screenshot of what is expected of the python call is shown below. At first, we thought this issue was because of caching in the OS, but quickly found out this was not the case after manually clearing the cache, as well as rebooting the system. After some further investigation, we found that Python3 was only getting the hostname on Clear Linux instead of the FQDN. Ciao relies on certain Python3 function calls. We reported this problem to the Clear Linux team as well, and they have since shipped a fix.

After running into these two issues, we spoke with the Ciao team about our setup, as we wanted to quickly iron out any other testing environment issues. We found out that the Ciao dev team usually uses Ubuntu as its development environment. With that in mind, we decided to move to an Ubuntu development environment to make it as consistent as possible. We have since then installed Ubuntu 16.10 on all the machines. Also at the recommendation of the Ciao team, we started using Ciao-Down.

Ciao-down is a single virtual machine development environment used by the Ciao development team to work on code locally, and simulate a full cluster environment. This is quite a boon for the team as far as independent development. It allows ease of setup, and quick deployment of changes to a simulated environment for testing. An example of how easy it is to set up:[1]

Listing 2. Simplicity of Ciao-Down setup

```
$ go get github.com/01org/ciao/testutil/ciao-down
$ ciao-down prepare
```

That's how simple it is to get a single environment setup working with Ciao-Down. Running `ciao-down prepare` spins up a new virtual machine on the local machine. It then installs any necessary tools needed, such as the Go programming language, an open-ssh server, and other tools. After `ciao-down prepare` is run, simply running `ciao-down connect` connects to the newly created virtual machine, which pulls from the Ciao repository on the local machine.

The only caveat to this was that the machine Ciao-Down is running on requires nested kernel-based virtualization. While this is not an uncommon feature, the steps for enabling it were a bit tedious. On the first attempt, we manually edited Grub's (the boot loader for Ubuntu) configuration in order to enable it. But after a helpful hint from one of the good folks at Intel, there is a simple one line command that enables it:

Listing 3. An easy way to enable nested KVM support

```
echo "options kvm-intel nested=1 | sudo tee /etc/modprobe.d/kvm-intel.conf"
```

With this set up, we were able to quickly get to work on the development of the first feature. While we are not currently using a fully deployed version of Ciao for development, we will work on getting Ciao tested with the new features in our full five NUC setup once progress has been made on development.

V. TEAM REVIEW

A. *Garrett Smith*

B. *Matthew Johnson*

VI. RETROSPECTIVE

| Week | Positives | Deltas | Actions |
|------|---|---|---|
| 1 | Learned about automated deployment for Ciao | Worked through several deployment issues | Attempted to deploy Ciao |
| 2 | Scripts to collect network statistics completed | Resolved old issues in deployment but uncovered more | Hardcoding variables in the ansible deployment resolved some new issues |
| 3 | Successfully initiated the cluster | Switched OS on the cluster to Ubuntu | Started working in single-vm mode |
| 4 | Wrote some of the module for OVS tunnels | OSU network DHCP will not assign IPs to unregistered MACs | Developed in ciao-down |
| 5 | OVS module mostly written, but not integrated | Discussed whether OVS bridges were required for OVS tunnels | Tried to set up local DHCP server on switch |
| 6 | Met with Manohar regarding OVS bridges | OVS bridges are required for OVS tunnels | Worked on progress report for midterm |
| 7 | | | |
| 8 | | | |
| 9 | | | |
| 10 | | | |

REFERENCES

- [1] m. mcastelino. (2017, January) Ciao single machine development and test environment. [Online]. Available: <https://github.com/01org/ciao/blob/master/DeveloperQuickStart.md>